

STATS 415 Lab7

Weijing Tang

Feb 22, 2018

1 Today's objectives

1. Randomly split dataset into training and test.
2. Practice how to implement Ridge regression and Lasso with λ chosen by cross-validation.
3. Report the training and test errors for the final models.
4. Compare the model performance among OLS, ridge regression and lasso.

2 Data and necessary packages

We will use the `Hitters` dataset from the `ISLR` package to explore two shrinkage methods: **ridge regression** and **lasso**. These are otherwise known as **penalized regression** methods.

```
library(ISLR)
data(Hitters)
```

We will use the `glmnet` package in order to perform ridge regression and the lasso. The main function in this package is `glmnet()`, which can be used to fit ridge regression models, lasso models, and more. This function has slightly different syntax from other model-fitting functions that we have encountered thus far in this book. In particular, we must pass in an `x` matrix as well as a `y` vector, and we do not use the `y ~ x` syntax. We will now perform ridge regression and the lasso in order to predict Salary on the `Hitters` data. Before proceeding ensure that the missing values have been removed from the data.

Recall that this dataset has some missing data in the response `Salary`. We use the `na.omit()` function to clean the dataset.

```
sum(is.na(Hitters))
```

```
## [1] 59
```

```
sum(is.na(Hitters$Salary))
```

```
## [1] 59
```

```
Hitters = na.omit(Hitters)
sum(is.na(Hitters))
```

```
## [1] 0
```

The `model.matrix()` function is particularly useful for creating `x`; not only does it produce a matrix corresponding to the 19 predictors but it also **automatically transforms any qualitative variables into dummy variables**. The latter property is important because `glmnet()` can only take numerical, quantitative inputs. It also returns the intercept term as the first column, so we need remove the first column by `[, -1]`.

```
library(glmnet)
X = model.matrix(Salary ~ ., Hitters)[, -1]  # don't want x to include the intercept
y = Hitters$Salary
```

3 Test-train Split

We begin by splitting the observations into a training set and a test set. We do this by creating a random vector, `train`, of elements equal to `TRUE` if the corresponding observation is in the training set, and `FALSE` otherwise. The vector `test` has a `TRUE` if the observation is in the test set, and a `FALSE` otherwise.

```
set.seed(1)  # if true, belong to train data, if false, test data
train=sample(c(TRUE,FALSE), nrow(Hitters),rep=TRUE)
test=(!train)
```

4 Ridge Regression

We first illustrate **ridge regression**, which can be fit using `glmnet()` with `alpha = 0` and seeks to minimize

$$\sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2.$$

Question: What is the model if $\lambda = 0$?

more variances included;
lambda=inf no variances included

Notice that the intercept is **not** penalized. Also, note that that ridge regression is **not** scale invariant like the usual unpenalized regression. Thankfully, `glmnet()` takes care of this internally. **It automatically standardizes predictors for fitting, then reports fitted coefficient using the original scale.**

Some commonly used arguments of `glmnet()`: 1. **x**: input matrix 2. **y**: response variable 3. **lambda**: the `glmnet()` function performs ridge regression for an automatically selected range of λ values by default. We can also set a grid of values. Here we choose to implement the function over a grid of values ranging from $\lambda = 10^{10}$ to $\lambda = 10^{-2}$, essentially covering the full range of scenarios from the null model containing only the intercept, to the least squares fit. 4. **alpha**: `alpha=0` the ridge penalty and `alpha=1` the lasso penalty. 5. **standardize**: the `glmnet()` function standardizes the variables so hat they are on the same scale by default. To turn off this default setting, use the argument `standardize=FALSE`. 6. **thresh**: convergence threshold for optimization method.

```
grid=10^seq(10,-2,length=100)
ridge.mod=glmnet(X[train,],y[train],alpha=0,lambda=grid)
```

Associated with each value of λ is a vector of ridge regression coefficients, stored in a matrix that can be accessed by `coef()`. In this case, it is a 20×100 matrix, with 20 rows (one for each predictor, plus an intercept) and 100 columns (one for each value of λ). **19+1intercept**


```
dim(coef(ridge.mod))
```

```
## [1] 20 100
```

Question: What do you expect the coefficient estimates to be if we increase λ ? Why?

We can support our conclusion by the following two plots.

```
par(mfrow = c(1, 2))
plot(ridge.mod)
plot(ridge.mod, xvar = "lambda", label = TRUE)
```



Question: Is there any coefficient is forced to zero?

Here we can again use `predict()` function to make predictions for `newx` and a given λ . We can also obtain the ridge regression coefficients for a new value of λ , say 50:

```
predict(ridge.mod, s=50, type="coefficients")[1:20,]
```


##	(Intercept)	AtBat	Hits	HmRun	Runs
##	-144.43422527	-0.02343557	2.81521875	-0.54884891	1.59947732
##	RBI	Walks	Years	CAtBat	CHits
##	1.30877723	0.65795601	-0.11157223	0.01693560	0.15314382
##	CHmRun	CRuns	CRBI	CWalks	LeagueN
##	0.02074275	0.16694273	0.12683534	-0.25530284	112.48936029
##	DivisionW	PutOuts	Assists	Errors	NewLeagueN
##	-59.10973952	0.18318693	-0.26560638	-2.03609898	-25.31027138

4.1 Choose λ by cross-validation

λ is tuning parameter here. We use cross-validation to select a good λ value, which should be based only on training data.

We can do this using the built-in cross-validation function, `cv.glmnet()`. By default, the function performs ten-fold cross-validation, though this can be changed using the argument `nfolds`. Note that we set a random seed first so our results will be reproducible, since the choice of the cross-validation folds is random.

```
set.seed(1)
cv.out=cv.glmnet(X[train,],y[train],alpha=0)
plot(cv.out)
```



The plot illustrates the MSE for the λ s considered. Two lines are drawn. The first is the λ that gives the smallest MSE. The second is the λ that gives an MSE within one standard error of the smallest.

```
bestlam=cv.out$lambda.min
bestlam
```

```
## [1] 25.00932
```

Therefore, we see that the value of λ that results in the smallest crossvalidation error is 25.

4.2 Calculating training and test error

```
# training MSE
ridge.pred_train=predict(ridge.mod,s=bestlam,newx=X[train,])
mean((ridge.pred_train-y[train])^2)
```

```
## [1] 67968.27
```

The training MSE is 67968.27.

```
# test MSE
ridge.pred_test=predict(ridge.mod,s=bestlam,newx=X[test,])
mean((ridge.pred_test-y[test])^2)
```

```
## [1] 154665.8
```

The test MSE is 154665.8.


5 Lasso

We now illustrate **lasso**, which can be fit using `glmnet()` with `alpha = 1` and seeks to minimize

$$\sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p |\beta_j|.$$

Like ridge, lasso is not scale invariant. The only difference between ridge regression and lasso is the penalty. Other than the change `alpha=1`, we proceed just as we did in fitting ridge regression model.


```
lasso.mod=glmnet(X[train,],y[train],alpha=1,lambda=grid)
par(mfrow = c(1, 2))
plot(lasso.mod)
plot(lasso.mod, xvar = "lambda", label = TRUE)
```



Question: Is there any coefficient is forced to zero? What if compared with ridge regression?

We now perform cross-validation and compute the associated test error.

```
set.seed(1)
cv.out=cv.glmnet(X[train,],y[train],alpha=1)
plot(cv.out)
```



```
# best lambda
bestlam=cv.out$lambda.min
bestlam
```

```
## [1] 10.57705
```

```
# training error
lasso.pred_train=predict(lasso.mod,s=bestlam,newx=X[train,])
mean((lasso.pred_train-y[train])^2)
```

```
## [1] 71678.76
```

```
# test error
lasso.pred_test=predict(lasso.mod,s=bestlam,newx=X[test,])
mean((lasso.pred_test-y[test])^2)
```

```
## [1] 161022.6
```

The test MSE of lasso is larger than that of ridge regression. However, the lasso has a substantial advantage over ridge regression in that the resulting coefficient estimates are sparse.

```
lasso.coef=predict(lasso.mod,type="coefficients",s=bestlam)[1:20,]
lasso.coef
```

```
## (Intercept)      AtBat      Hits      HmRun      Runs
## -1.753112e+02  0.000000e+00  4.207687e+00  0.000000e+00  0.000000e+00
##           RBI      Walks      Years      CAtBat      CHits
```

```
## 7.704335e-01 0.000000e+00 0.000000e+00 0.000000e+00 2.685458e-01
##          CHmRun          CRuns          CRBI          CWalks          LeagueN
## 0.000000e+00 0.000000e+00 0.000000e+00 -2.526989e-03 7.425729e+01
##      DivisionW      PutOuts      Assists      Errors      NewLeagueN
## -4.324510e+01 1.632754e-01 -3.591127e-01 -1.949470e-02 0.000000e+00
```

Here we see that 10 of 19 coefficient estimates are exactly zero. So the lasso chosen by cross-validation contains only 9 variables.

6 Compare the model performance

We fit the ordinary least square(null model) and calculate training and test error.

```
lm.mod = lm(Salary~.,data = Hitters[train,])
# training error
lm.pred_train = predict(lm.mod,Hitters[train,])
mean((y[train]-lm.pred_train)^2)
```

```
## [1] 61988.36
```

```
# test error
lm.pred_test = predict(lm.mod,Hitters[test,])
mean((y[test]-lm.pred_test)^2)
```

```
## [1] 160105.6
```

Model	Train Error	Test Error
OLS	61988.36	160105.6
Ridge Regression	67968.27	154665.8
Lasso	71678.76	161022.6

Question: We can see that OLS has the lowest training Error. Does it always happen?

Question: Which model should we choose if we want a model to do better at prediction?

Question: Which model should we choose if we prefer a model simple to interpret?

7 Resources and References

1. <https://davidalpiaz.github.io/r4sl/regularization.html#ridge-regression>